

Introduction

Machine Learning is the study of algorithms that can learn from experience, typically in the form of observational data or interactions with an environment, its performance improves.

Core components that follow us around, no matter what kind of machine learning we tackle:

1. The *data* that we can learn from
2. A *model* of how to transform the data
3. An *objective function* that quantifies how well (or badly) the model is doing
4. An *algorithm* to adjust the model's parameters to optimize the objective function

Note 0.1

Data

Each data point/instance consists of a set of attributes called *features* (or co-varies/inputs)/

In supervised learning, our goal is to predict the value of a special attribute, called the *label/target* which is not part of the model's input.

In cases when every example is characterized by the same number of numerical features, we say that the inputs are fixed-length vectors and we call the constant length of the vectors the *dimensionality* of the data. However, often we have to work with varying lengths of data.

By **model**, we denote the computational machinery for ingesting data of one type, and spitting out predictions of a possibly different type. Specifically, we are interested in *statistical models* that can be estimated from data.

Deep learning is differentiated from the classical approaches principally by the set of power models that it focuses on. We often chain many successive transformations of the data from top to bottom.

Note 0.2

In order to develop a formal mathematical system for learning machines, we need to have objective functions to see how good or bad our models are. We often define objective functions so that lower is better by convention, which is why we call them *loss functions*.

When trying to predict numerical values, we can often use *squared error* as a loss function, but for classification, the most common objective function is to minimize error rate (% of wrong predictions). **During optimization, we think of the loss as a function of the model's parameters, and treat the training dataset as a constant.** We will learn the best values of our model's parameters by minimizing the loss incurred on a set consisting of some number of examples collected for training.


We typically want to split the available data into a *training dataset* and a *test set* which is held for evaluation. When a model performs well on the training set but fails to generalize unseen data, we say that it is **overfitting** on training data.

Now, once we have a data source, a model, and a well defined objective function, we need an algorithm capable of searching for the best possible parameters for minimizing the loss function. Popular optimization algorithms for deep learning are based on an approach known as **gradient descent**

Gradient Descent: an optimization algorithm where the method check to see (at each step) for each parameter, how that training set would change if you perturbed that parameter by just a small amount. It then updates the parameter in a direction that lowers the loss.

Neural Networks and other complicated models often have many hundreds or thousands of inputs, so using calculus/statistics to find the minimum is often not feasible. Instead, we start at any old input and find the slope of the function where you are, and shift towards the direction that minimizes the loss.

However, there are many local minimums we can land in (see below) and we might not actually get the most minimum.



This carries over to neural networks, and its really hard to find the absolute minimum.

Note: if you make your step sizes proportional to the slope, then as your slope flattens out your steps get smaller and smaller to help you from overshooting.

When we add more inputs, we instead want to use the idea of **gradient**, which is the direction of the steepest increase and also tells us just how steep it is. We can use this to move our parameters in the opposite direction of the steepest increase and know how much to move it.

Steps:

1. Compute ∇C
2. Small step in $-\nabla C$ direction
3. Repeat

∇C is a vector that tells us how to “nudge” all weights and biases to cause the most rapid decrease to the cost function. The sign of each value in the vector tells us which direction to move it.

Kinds of Machine Learning

Supervised learning describes tasks where we are given a dataset containing both features and labels and asked to produce a model that predicts the labels when given input features.

The supervision comes into play because for choosing parameters we produce the model with a dataset consisting of labeled examples. We are interested in estimating the conditional probability of a label given input features.

When labels take on arbitrary numerical values (even if its within some interval), we consider this to be a **regression problem**.

In most our chapters, we will try to learn models that minimize the squared error loss function.

In **classification**, we want our model to look at features and predict which category/class it belongs to.

Reinforcement learning gives a very general statement of a problem in which an agent interacts with an environment over a series of time steps. At each step, the agent receives some *observation* from the environment and must chose an *action* that is subsequently transmitted back into the environment leading to some reward by the ML model.